

Adapted from Computer Organization the Hardware/Software Interface - 5th

1

The Processo



Introduction

- CPU performance factors
 - Instruction count
 - Determined by ISA and compiler
 - CPI and Cycle time
 - Determined by CPU hardware
- We will examine two MIPS implementations
 - A simplified version
 - A more realistic pipelined version
- Simple subset, shows most aspects
 - Memory reference: 1w, sw
 - Arithmetic/logical: add, sub, and, or, slt
 - Control transfer: beq, j

2

2



Instruction Execution

- PC \rightarrow instruction memory, fetch instruction
- Register numbers → register file, read registers
- Depending on instruction class
 - Use ALU to calculate
 - · Arithmetic result
 - · Memory address for load/store
 - · Branch target address
 - Access data memory for load/store
 - PC ← target address or PC + 4

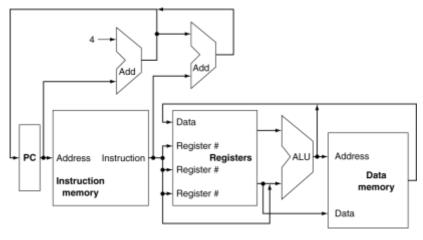
3

3

The Processo



CPU Overview



4

4



Execution Model

- Instruction fetch: PC → instruction address
- Instruction decode: register operands → register file
- Instruction execute:
 - Load/store: compute a memory address
 - Arithmetic: compute an arithmetic result
- Write back:
 - Load/store: store a value to a register or a memory location
 - Arithmetic: store a result of register file

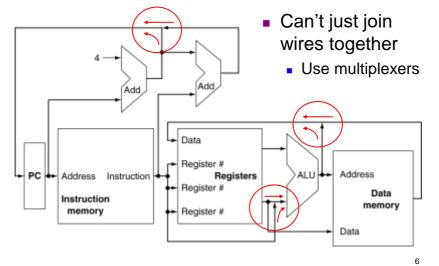
.

5

The Process



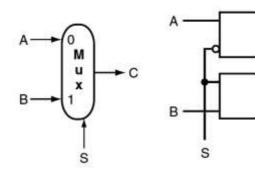
Multiplexers



6



Multiplexer



$$C = A\bar{S} + BS$$

7

7

The Processo

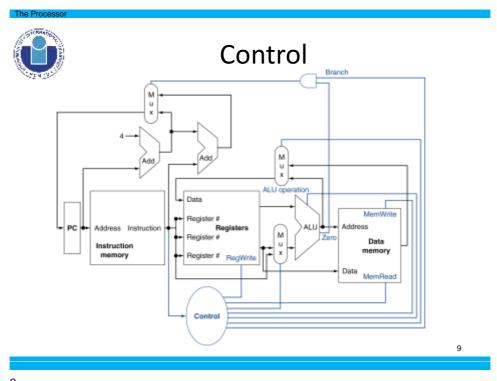


Control vs. Data signals

- Control signal: used for multiplexer selection or for directing the operation of a functional unit
- Data signal: contains information that is operated on by a functional unit

8

8



9





Logic Design Basics

- Information encoded in binary
 - Low voltage = 0, High voltage = 1
 - One wire per bit
 - Multi-bit data encoded on multi-wire buses
- Combinational element
 - Operate on data
 - Output is a function of input
- State (sequential) elements
 - Store information

10

10



Combinational Elements

AND-gate

$$- Y = A \& B$$

Multiplexer

$$- Y = S ? I1 : I0$$



Adder

$$-Y = A + B$$



• Arithmetic/Logic Unit

$$-Y = F(A,B)$$



1

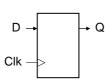
11

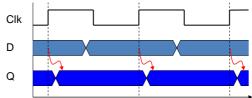
The Process



Sequential Elements

- Register: stores data in a circuit
 - Uses a clock signal to determine when to update the stored value
 - Edge-triggered: update when Clk changes from 0 to 1





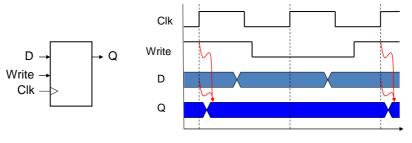
12

12



Sequential Elements

- Register with write control
 - Only updates on clock edge when write control input is 1
 - Used when stored value is required later



13

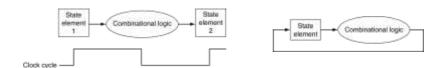
13

The Processo



Clocking Methodology

- Combinational logic transforms data during clock cycles
 - Between clock edges
 - Input from state elements (a memory or a register), output to state element
 - Longest delay determines clock period



14

14



Building a Datapath

- Datapath
 - Elements that process data and addresses in the CPU
 - Registers, ALUs, mux's, memories, ...
- We will build a MIPS datapath incrementally
 - Refining the overview design

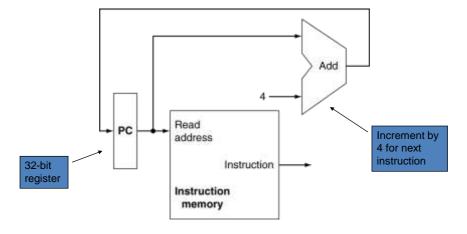
15

15

The Processo



Instruction Fetch



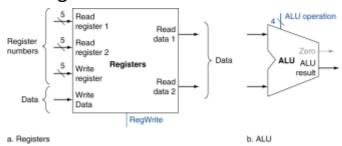
16

16



R-Format Instructions

- Ex. add, sub, and, or, slt
- Read two register operands
- · Perform arithmetic/logical operation
- Write register result



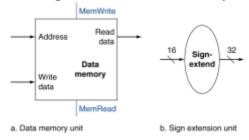
17

The Processo



Load/Store Instructions

- Read register operands
- Calculate address using 16-bit offset
 - Use ALU, but sign-extend offset
- Load: Read memory and update register
- Store: Write register value to memory



18



Branch Instructions

- Brach taken: condition is satisfied and the Program Counter (PC) register becomes the branch target
- Branch not taken: PC becomes the address of the next instruction
- Datapath:
 - Compute the branch target
 - Compare the registers

19

19

he Process



Branch Instructions

- Read register operands
- Compare operands
 - Use ALU, subtract and check Zero output
- Calculate target address
 - Sign-extend displacement
 - Shift left 2 places (word displacement)
 - Add to PC + 4
 - · Already calculated by instruction fetch

20

20

Branch Instructions PC+4 from instruction datapath Just Branch re-routes target wires Shift left 2 4 ALU operation register 1 Read data 1 Read register 2 To branch ALU Zero Registers control logic Write register Read

Sign-

Write

Reg/Write

21

21

The Processo



Composing the Elements

Sign-bit wire replicated

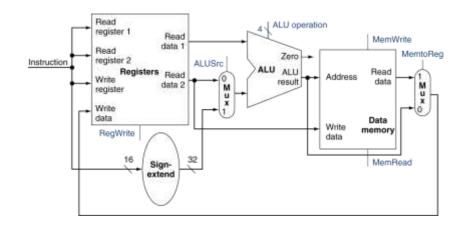
- First-cut data path does an instruction in one clock cycle
 - Each datapath element can only do one function at a time
 - Hence, we need separate instruction and data memories
- Use multiplexers where alternate data sources are used for different instructions
 - ALU input: register value/immediate
 - Register file write data: results from the ALU/Memory

22

22

Ü

R-Type/Load/Store Datapath



23

23

Full Datapath Add result left 2 Read register 1 MemWrite 1 - 1 Zem register 2 ALU ALU Instruction Write Regist Instruction register Write Data Write MemRead Sign-extend

24



ALU Control

ALU used for

Load/Store: F = addBranch: F = subtract

- R-type: F depends on funct field

ALU control	Function		
0000	AND		
0001	OR		
0010	add		
0110	subtract		
0111	set-on-less-than		
1100	NOR		

2

25

he Processo



ALU Control

• Assume 2-bit ALUOp derived from opcode

Combinational logic derives ALU control

Combinational logic actives / Les control									
opcode	ALUOp	Operation	funct	ALU function	ALU control				
lw	00	load word	XXXXXX	add	0010				
sw	00	store word	XXXXXX	add	0010				
beq	01	branch equal	XXXXXX	subtract	0110				
R-type	10	add	100000	add	0010				
		subtract	100010	subtract	0110				
		AND	100100	AND	0000				
		OR	100101	OR	0001				
		set-on-less-than	101010	set-on-less-than	0111				

26

26



Multiple Levels of Decoding

- Instruction opcode + function => ALUOp => ALU control lines
 - Smaller main controller
 - Faster main controller (potentially)

27

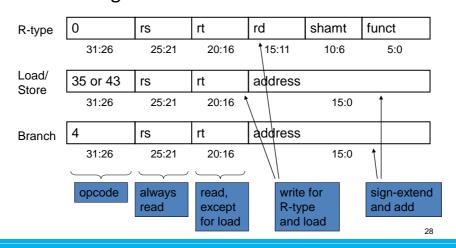
27

he Processo



The Main Control Unit

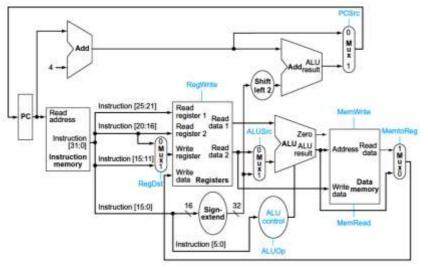
• Control signals derived from instruction



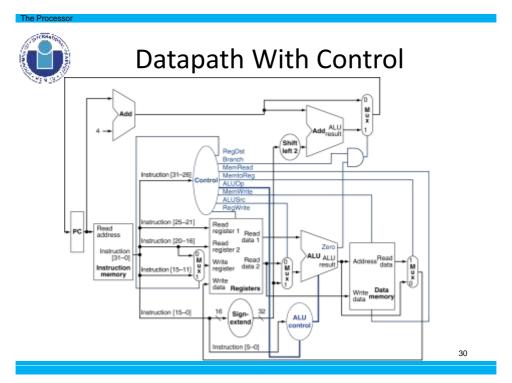
28



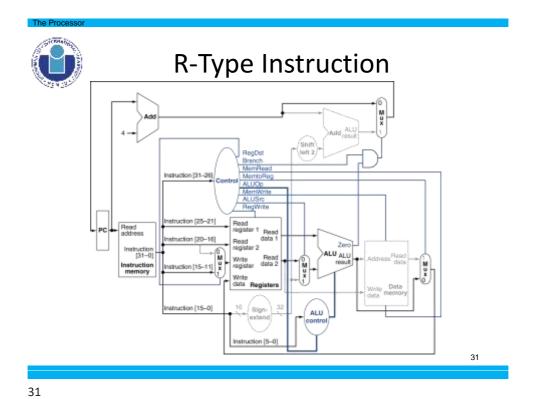
Datapath with Instruction Fields



29



30



Load Instruction

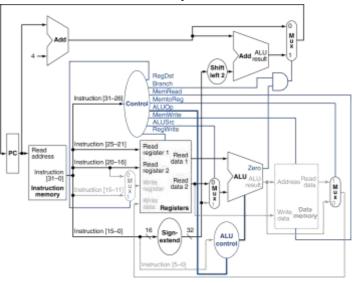
HegDa Shift Instruction [31-28] Control Membrand Membrand

HCMIU-SCSE 16

32

Ü

Branch-on-Equal Instruction



33



Implementing Jumps

- Jump uses word address
- Update PC with concatenation of
 - Top 4 bits of old PC
 - 26-bit jump address
 - -00
- Need an extra control signal decoded from opcode

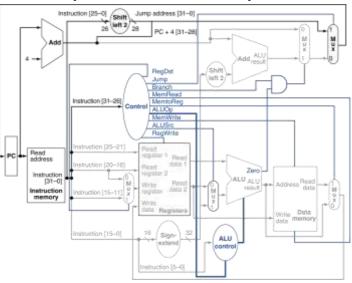
Jump 2 address 25:0

34

34



Datapath With Jumps Added



35

35

The Processo



Performance Issues

- Longest delay determines clock period
 - Critical path: load instruction
 - Instruction memory \rightarrow register file \rightarrow ALU \rightarrow data memory \rightarrow register file
- Not feasible to vary period for different instructions
- Violates design principle
 - Making the common case fast
- We will improve performance by pipelining

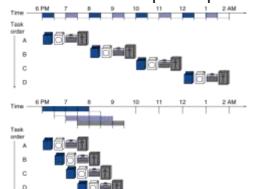
36

36



Pipelining Analogy

- Pipelined laundry: overlapping execution
 - Parallelism improves performance



- Four loads:
 - Speedup= 8/3.5 = 2.3
- Non-stop:
 - Speedup= 2n/0.5n + 1.5 ≈ 4= number of stages

37

37

The Process



MIPS Pipeline

- Five stages, one step per stage
 - 1. IF: Instruction fetch from memory
 - 2. ID: Instruction decode & register read
 - 3. EX: Execute operation or calculate address
 - 4. MEM: Access memory operand
 - 5. WB: Write result back to register

38

38



Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages

Compare pipelined datapath with single-cycle

datapath

datapath							
Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time	
lw	200ps	100 ps	200ps	200ps	100 ps	800ps	
sw	200ps	100 ps	200ps	200ps		700ps	
R-format	200ps	100 ps	200ps		100 ps	600ps	
beq	200ps	100 ps	200ps			500ps	

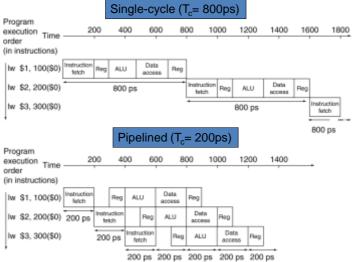
39

39

he Processo



Pipeline Performance



40

40



Pipeline Speedup

- If all stages are balanced
 - i.e., all take the same time

 $\label{eq:time-potential} Time\ between\ instruction_{nonpipelined} = \frac{Time\ between\ instruction_{nonpipelined}}{Number\ of\ pipe\ stages}$

- If not balanced, speedup is less
- Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease

4

41

he Processo



Pipelining and ISA Design

- MIPS ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle

42

42



Hazards

- Situations that prevent starting the next instruction in the next cycle
- Structure hazards
 - A required resource is busy
- Data hazard
 - Need to wait for previous instruction to complete its data read/write
- Control hazard
 - Deciding on control action depends on previous instruction

43

43

he Processo



Structure Hazards

- Conflict for use of a resource
- In MIPS pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to stall for that cycle
 - Would cause a pipeline "bubble"
- Hence, pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches

44

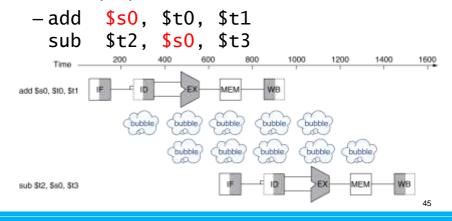
44





Data Hazards

An instruction depends on completion of data access by a previous instruction

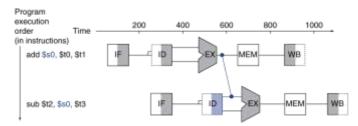


45



Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath

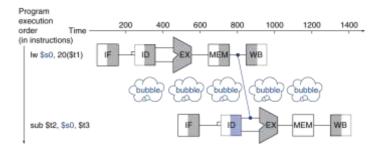


46



Load-Use Data Hazard

- · Can't always avoid stalls by forwarding
 - If value not computed when needed
 - Can't forward backward in time!



47

47

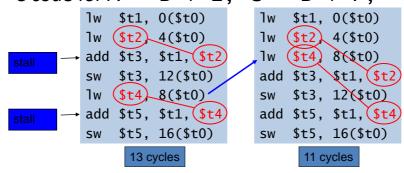
he Processo



Code Scheduling to Avoid Stalls

 Reorder code to avoid use of load result in the next instruction

• C code for A = B + E; C = B + F;



48

48



Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - · Still working on ID stage of branch
- In MIPS pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage

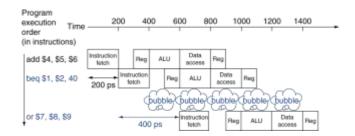
49

49



Stall on Branch

 Wait until branch outcome determined before fetching next instruction



50

50



Branch Prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In MIPS pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

5

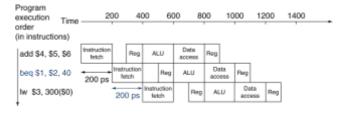
51

The Processo

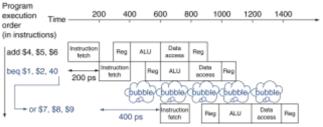


MIPS with Predict Not Taken

Prediction correct



Prediction incorrect



52

52



More-Realistic Branch Prediction

- Static branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - · Predict backward branches taken
 - · Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - · e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - · When wrong, stall while re-fetching, and update history

53

53

The Processo



Pipeline Summary

The BIG Picture

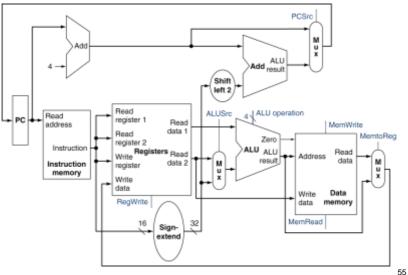
- Pipelining improves performance by increasing instruction throughput
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

54

54



Single-cycle Datapath



55

The Process

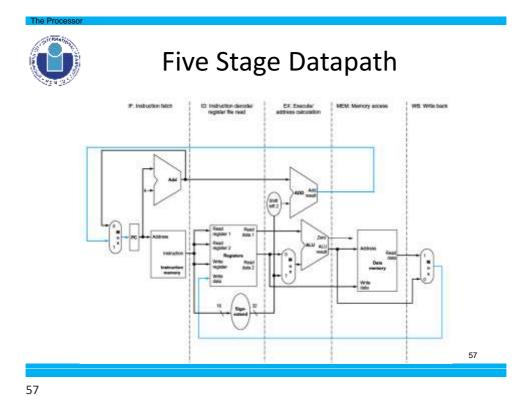


(Multi-cycle) Pipelined Datapath

- Up to five instruction will be in execution in one clock cycle
- Separate the datapath into five pieces:
 - IF: Instruction fetch
 - ID: Instruction decode and register file read
 - EX: Execution and address calculation
 - MEM: Data memory access
 - WB: Write back

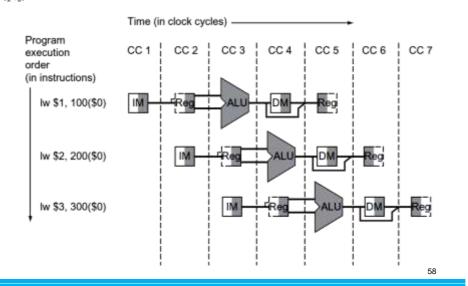
56

56



The Proce

Instructions Execution

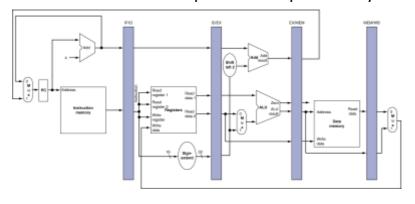


58



Pipeline registers

- Need registers between stages
 - To hold information produced in previous cycle



59

59

he Processo

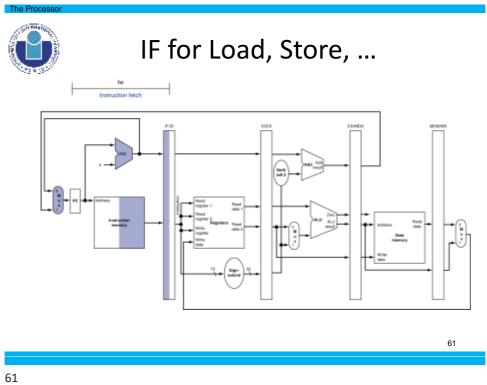


Pipeline Operation

- Cycle-by-cycle flow of instructions through the pipelined datapath
 - "Single-clock-cycle" pipeline diagram
 - Shows pipeline usage in a single cycle
 - Highlight resources used
 - c.f. "multi-clock-cycle" diagram
 - Graph of operation over time
- We'll look at "single-clock-cycle" diagrams for load & store

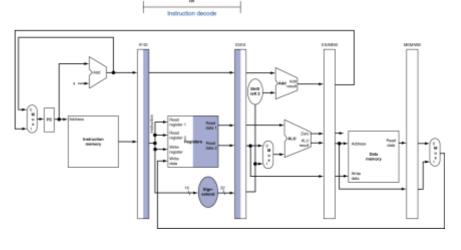
60

60

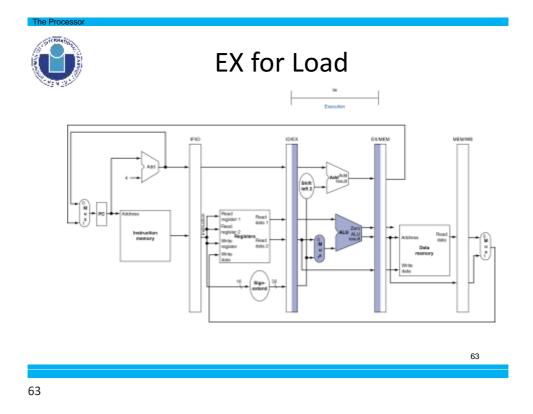


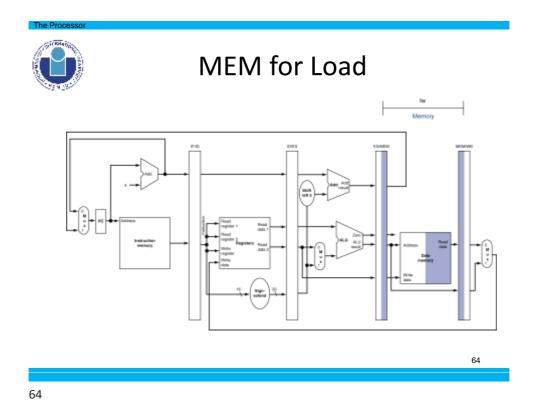


ID for Load, Store, ...



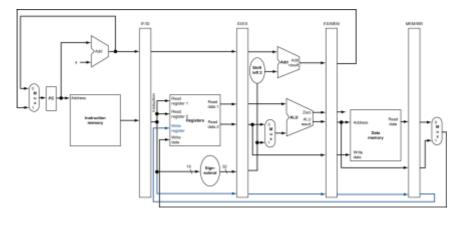
62





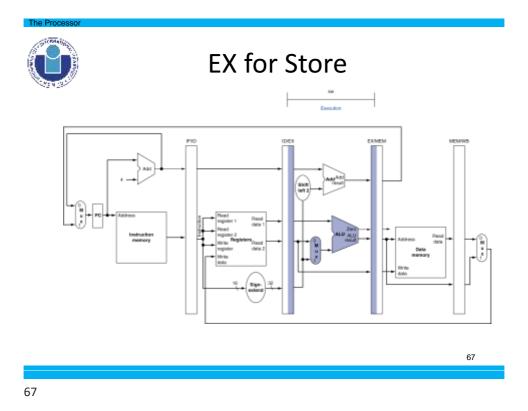
WB for Load Wrong register number 65

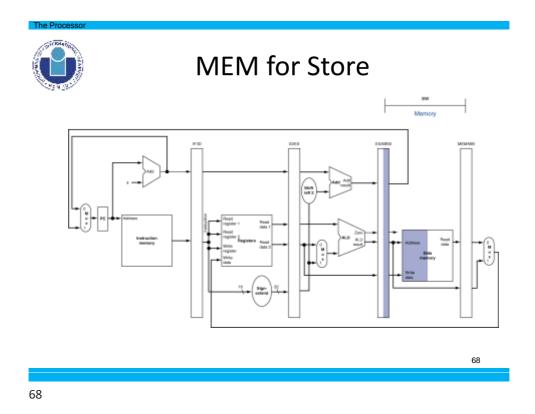
Corrected Datapath for Load

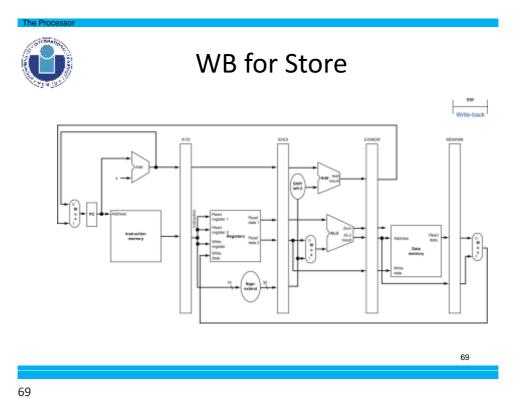


66

66

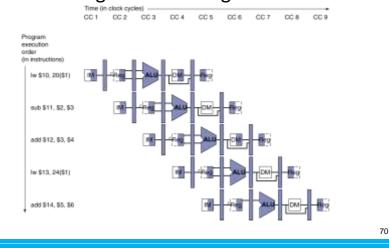






Multi-Cycle Pipeline Diagram

• Form showing resource usage

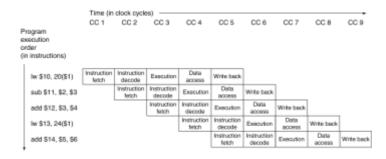


70



Multi-Cycle Pipeline Diagram

Traditional form

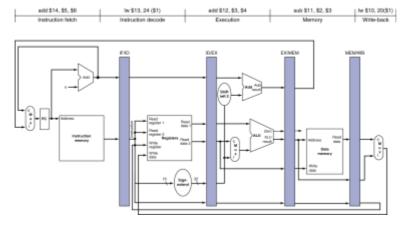


71



Single-Cycle Pipeline Diagram

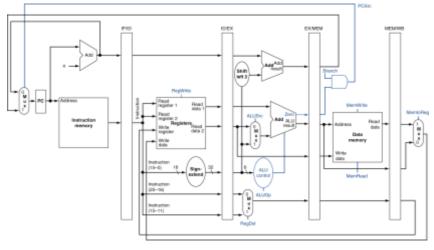
State of pipeline in a given cycle



72



Pipelined Control (Simplified)

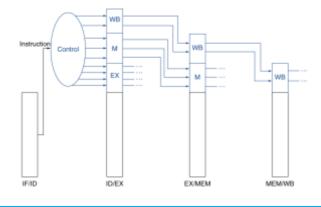


73

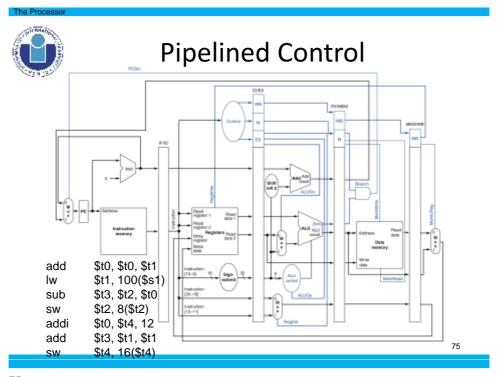


Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation



74



75



Data Hazards in ALU Instructions

• Consider this sequence:

```
sub $2, $1,$3
and $12,$2,$5
or $13,$6,$2
add $14,$2,$2
sw $15,100($2)
```

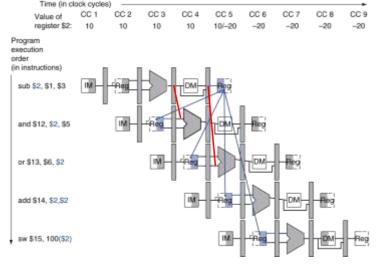
- · We can resolve hazards with forwarding
 - How do we detect when to forward?

76

76



Dependencies & Forwarding



77



Detecting the Need to Forward

- · Pass register numbers along pipeline
 - e.g., ID/EX.RegisterRs = register number for Rs sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
 - ID/EX.RegisterRs, ID/EX.RegisterRt
- Data hazards when
 - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
 - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
 - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
 - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

Fwd from EX/MEM pipeline reg

Fwd from MEM/WB pipeline reg

78

78



Detecting the Need to Forward

- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not \$zero
 - EX/MEM.RegisterRd ≠ 0,MEM/WB.RegisterRd ≠ 0

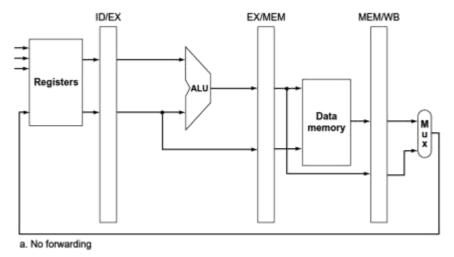
79

79

The Processo



No Forwarding Paths

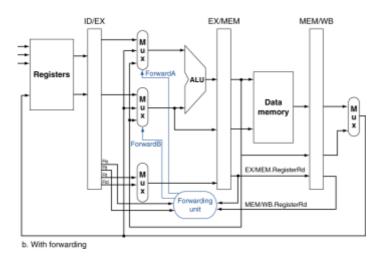


80

80



Forwarding Paths



81

81

The Processo



Forwarding Control Values

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

82

82



Forwarding Conditions

- EX hazard
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRs)) ForwardA = 10
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRt)) ForwardB = 10
- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRs)) ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRt)) ForwardB = 01

8

83

he Processo



Double Data Hazard

• Consider the sequence:

add \$1,\$1,\$2 add \$1,\$1,\$3 add \$1,\$1,\$4

- · Both hazards occur
 - Want to use the most recent
- Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true

84

84



Revised Forwarding Condition

MEM hazard

ForwardB = 01

- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
 and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
 and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
 and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
 ForwardA = 01
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
 and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
 and (EX/MEM.RegisterRd = ID/EX.RegisterRt))
 and (MEM/WB.RegisterRd = ID/EX.RegisterRt))

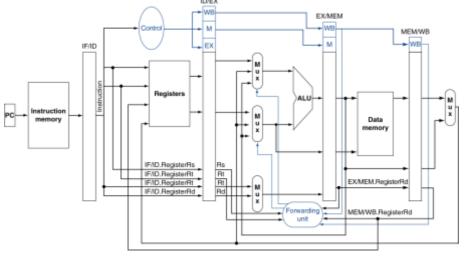
8.5

85

he Processo



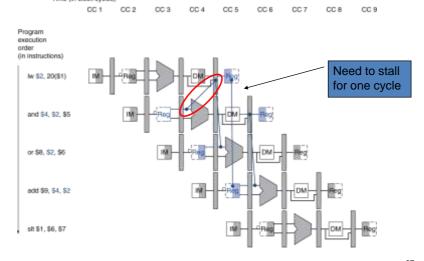
Datapath with Forwarding



86



Load-Use Data Hazard



87



Load-Use Hazard Detection

- Check when using instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
 - IF/ID.RegisterRs, IF/ID.RegisterRt
- Load-use hazard when
 - ID/EX.MemRead and ((ID/EX.RegisterRt = IF/ID.RegisterRs) or (ID/EX.RegisterRt = IF/ID.RegisterRt))
- If detected, stall and insert bubble

88

88



How to Stall the Pipeline

- Force control values in ID/EX register to 0
 - EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
 - Using instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for \(\frac{1}{W} \)
 - Can subsequently forward to EX stage

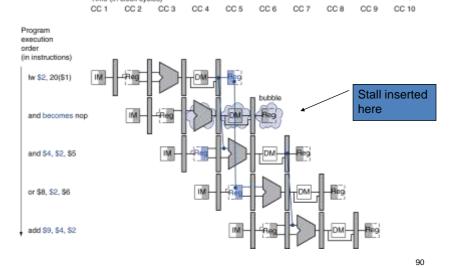
89

89

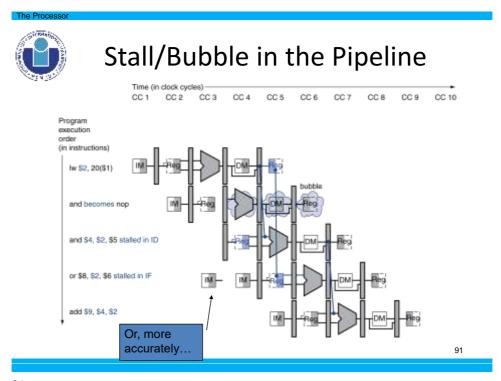
The Processo



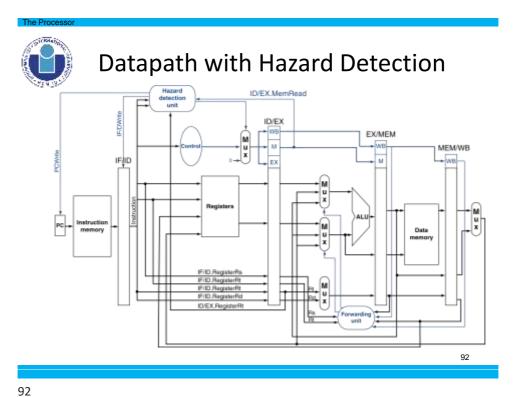
Stall/Bubble in the Pipeline



90



91



92





Stalls and Performance

The BIG Picture

- Stalls reduce performance
 - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

93

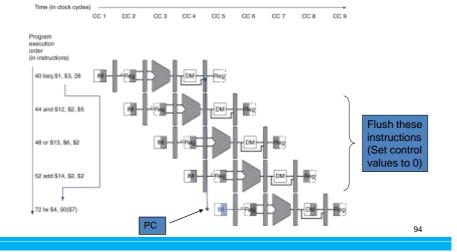
93

The Processo



Branch Hazards

If branch outcome determined in MEM



94



Reducing Branch Delay

- Move hardware to determine outcome to ID stage
 - Target address adder
 - Register comparator
- Example: branch taken

```
36:
    sub
         $10, $4, $8
40:
    beq
          $1, $3, 7
          $12, $2, $5
44:
    and
          $13, $2, $6
48:
    or
52:
    add $14, $4, $2
         $15, $6, $7
56:
    slt
```

72: lw \$4, 50(\$7)

9

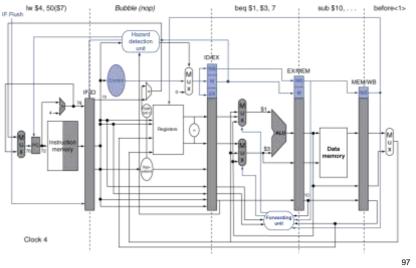
95

Example: Branch Taken and \$12, \$2, \$5 beq \$1, \$3, 7 sub \$10, \$4, \$8 before-1> Clock 3

96



Example: Branch Taken



97



Data Hazards for Branches

 If a comparison register is a destination of 2nd or 3rd preceding ALU instruction

• Can resolve using forwarding

98

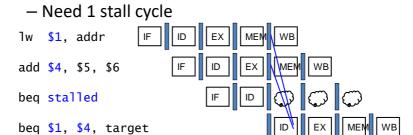
98

THE FIOLESS



Data Hazards for Branches

 If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction



99

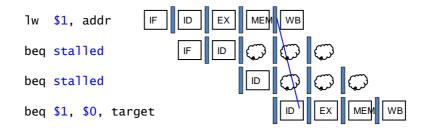
99

he Processo



Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles



100

100



Dynamic Branch Prediction

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - · Check table, expect the same outcome
 - · Start fetching from fall-through or target
 - · If wrong, flush pipeline and flip prediction

101

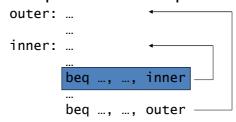
101

The Process



1-Bit Predictor: Shortcoming

Inner loop branches mispredicted twice!



- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of outer loop next time around

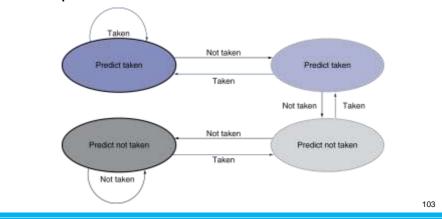
102

102



2-Bit Predictor

 Only change prediction on two successive mispredictions



103

The Processo



Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- · Branch target buffer
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately

104

104